

Practical FaaS Performance Variability Implications in the Serverless Image Optimization Scenario

Master's Thesis

Author

Piotr Witkowski

414360

witkowski@campus.tu-berlin.de

Advisor

Trever Schirmer

Examiners

Prof. Dr.-Ing. David Bermbach

Prof. Dr. habil. Odej Kao

Technische Universität Berlin, 2025

Fakultät Elektrotechnik und Informatik

Fachgebiet Scalable Software Systems

Practical FaaS Performance Variability Implications in the Serverless Image Optimization Scenario

Master's Thesis

Submitted by:
Piotr Witkowski
414360
witkowski@campus.tu-berlin.de

Technische Universität Berlin
Fakultät Elektrotechnik und Informatik
Fachgebiet Scalable Software Systems

2025

Hiermit versichere ich, dass ich die vorliegende Arbeit eigenständig ohne Hilfe Dritter und ausschließlich unter Verwendung der aufgeführten Quellen und Hilfsmittel angefertigt habe. Alle Stellen die den benutzten Quellen und Hilfsmitteln unverändert oder sinngemäß entnommen sind, habe ich als solche kenntlich gemacht.

Sofern generische KI-Tools verwendet wurden, habe ich Produktnamen, Hersteller, die jeweils verwendete Softwareversion und die jeweiligen Einsatzzwecke (z. B. sprachliche Überprüfung und Verbesserung der Texte, systematische Recherche) benannt. Ich verantworte die Auswahl, die Übernahme und sämtliche Ergebnisse des von mir verwendeten KI-generierten Outputs vollumfänglich selbst.

Die Satzung zur Sicherung guter wissenschaftlicher Praxis an der TU Berlin vom 8. März 2017* habe ich zur Kenntnis genommen.

Ich erkläre weiterhin, dass ich die Arbeit in gleicher oder ähnlicher Form noch keiner anderen Prüfungsbehörde vorgelegt habe.

(Unterschrift) Piotr Witkowski, Berlin, 20. Januar 2025

*https://www.static.tu.berlin/fileadmin/www/10000060/FSC/Promotion__Habilitation/Dokumente/Grundsatzgute_wissenschaftliche_Praxis_2017.pdf

Abstract

Kurzfassung

Contents

1	Introduction	8
1.1	Objectives of the Thesis	8
1.2	Structure	9
2	Background	10
2.1	Imaging and Photography	10
2.1.1	Digital imaging	11
2.1.2	Data and image compression	12
2.1.3	Raster and vector graphics	14
2.1.4	Raster image formats	15
2.2	Cloud Computing	17
2.2.1	Cloud deployment models	17
2.2.2	Event-driven architectures	17
2.2.3	Cloud service benchmarking	17
2.2.4	Performance variability of serverless platforms	17
2.3	Image Optimization for the Web	17
2.3.1	Impact of images on browsing experience	17
2.3.2	Bandwidth usage, visual quality, and browser support evolution	17
2.3.3	Optimizing latency of image delivery	17
2.3.4	Deployment methods for cloud-based image optimization	17
3	System Design	18
3.1	Image processing workload design	18
3.1.1	Serverless environment selection	18
3.1.2	Image manipulation library selection	18
3.1.3	Input image and output image selection	18
3.2	Benchmark design	18
3.2.1	Serverless platform and underlying hardware	18
3.2.2	Time of execution	18
3.2.3	Output formats and image widths	18
3.2.4	Available CPU / memory	18
3.2.5	Cold starts	18
3.3	Workload generation and benchmark execution	18
3.3.1	Serverless benchmark deployment	18
3.3.2	Scheduling workload executions	18
3.3.3	Results collection and data analysis	18
4	Results	19
4.1	Performance metrics	19
4.1.1	Cold start duration	19
4.1.2	Image processing duration	19
4.1.3	Memory utilization	19
4.2	Performance variability	19
4.2.1	General variability in performance metrics	19
4.2.2	Temporal variability	19
4.2.3	By image format	19
4.2.4	By platform configuration	19

5	Discussion	20
5.1	Performance variability analysis	20
5.2	User vs operator perspective	20
5.3	Cloud computing as a utility	20
6	Related Work	21
6.1	Serverless performance research	21
6.2	Image processing research	21
7	Conclusion	22

1 Introduction

The digital revolution has fundamentally transformed the way we create, consume, and interact with visual information. Images are no longer only physical items around us, they are inherent components of our online experience, shaping perceptions and influencing decisions. As the Internet continues to evolve, so too grows the demand for rich and engaging visual content. This demand, however, comes at a cost. Images, while visually appealing, are significantly impacting website performance, not only by generating the highest number of requests, but also by consuming the most bytes during page load of an average site. Existing studies have shown a strong correlation between page load times and user engagement, conversion rates, and even search engine rankings. Minimizing image size is therefore fundamental for efficient bandwidth use and plays a crucial role in optimizing website performance.

Image optimization practices for the web have evolved over time, following developments of the state-of-the-art image formats and codecs. These solutions perform operations such as image resizing and format conversion, and may allow for storing optimized versions of original files for future use. While the optimization itself is a relatively specialized task, its deployment models follow familiar and established patterns in the general software industry: we can find them in software-as-a-service offerings, integrated into web hosting and content delivery platforms, but also as open-source solutions for self-hosting.

This last option, leveraging open-source image processing libraries, presents a compelling opportunity for our research. Unlike proprietary solutions, open-source projects offer the flexibility to be deployed across a variety of environments, granting builders granular control over the optimization process. This allows not only to measure influence of different libraries and optimization parameters, but also – and crucially for this thesis – enables realistic benchmarking of the underlying platforms on which these operations are performed. By employing these workloads, our research aims to accurately measure cloud service performance under the realistic conditions of real-world image optimization scenarios.

Given the event-driven nature of the optimization workload in response to requests for images, serverless cloud compute emerges as a well-suited deployment option for these systems. However, the varying demands for image processing, including different image sizes and diverse range of image formats present on the web, raise questions about the suitability and performance consistency of these platforms for the use-case. Furthermore, the emergence of function as a service (FaaS) platforms introduces a new dimension to cloud-based, self-hosted image optimization. FaaS offers a compelling execution environment for image processing tasks thanks to its on-demand scalability and event-driven nature. However, the performance variability inherent in FaaS platforms, attributed to diverse factors, may impact efficiency of the process and applicability of the environment for the use case.

1.1 Objectives of the Thesis

This thesis aims to investigate whether FaaS performance variability has measurable and meaningful influence for image optimization workloads. We will define set of representative metrics related to image optimization, reflecting realistic web usage. We will design and implement the optimization workflows for selected serverless environments. We will design and implement scheduling mechanism allowing us to simulate requests for optimized images. We will collect performance metrics across all parameters, apply statistical

methods for result comparisons, and visualize the results. We will investigate if more intensive optimization can be cost-efficient, considering optimization result caching, and content reuse for multiple viewers. We will provide recommendations that are within the scope of the benchmark, and answer the primary research question in the context of the scenario.

1.2 Structure

2 Background

This chapter provides readers with foundations for understanding the challenges and opportunities of serverless image processing. We begin by revisiting fundamental principles of imaging and photography. Subsequently, we describe core digital imaging concepts, emphasizing aspects relevant to web performance. We then examine the characteristics of serverless cloud computing, focusing on factors contributing to performance variability. Finally, we explore image optimization for the web, reviewing established techniques and the landscape of cloud-based solutions.

2.1 Imaging and Photography

Imaging is a broad term describing any process that helps us in visualizing and understanding physical objects, abstract concepts, or even sets of data. While this may be a purely philosophical activity, for the purpose of this thesis, we will define imaging as creating a visual representation - an image - of physical objects and scenes. Humans are able to perceive and interpret these images thanks to our visual system, that has evolved over millions of years. Developed from the simple light-sensitive cells of early organisms to the sophisticated eyes of modern humans, this system is providing us with the ability to see color, depth, or movement.

From the earliest cave paintings to modern photography and film, humans have been exploring ways to capture and recreate the visual world. Over centuries, scholars and artists showed great interest in the process of creating visual representations to understand our own mind and senses, as well as to be able to depict the world around us more accurately. The natural phenomenon of light passing through a small aperture, called camera obscura, helped us understand fundamental principles of optics and image formation, and according to some, greatly contributed to realism of European art during Middle Ages. We've been refining tools and techniques to enhance our vision, and improve everyday life. Devices like telescopes and microscopes allowed us to observe objects at different scales, unveiling worlds that would otherwise remain invisible to the naked eye. Eyeglasses and contact lenses correct refractive errors, demonstrating that our visual experiences can not only be externally altered, but also subjectively improved, for millions of people.

Photography emerged in the 19th century as a groundbreaking invention, offering a novel approach to image creation. Unlike paintings or drawings, which rely on an artist's interpretation, photography directly captures and permanently records the physical world using light-sensitive materials. Early photographic methods required long exposure times, often making it possible to photograph only still objects and posed portraits. The daguerreotype, which used a silver-coated copper plate, was one of the first widely adopted photographic processes. Later, more versatile film technologies significantly reduced exposure times and made photography more accessible to the wider public. This ability to easily and accurately record the physical world had significant consequences. Newspapers, now able to include actual photographs, found a powerful new way to share news and social commentary through photojournalism. On a personal level, photography allowed people to preserve their own histories through portraits and the documentation of special occasions.

Since the beginning of the 21st century, digital photography has become the dominant form of capturing and sharing images. Initially limited to standalone devices, digital cameras are now integrated into smartphones and personal computers. Electronic sensors capture and

convert light into digital data, facilitating easy manipulation and image sharing, including over the Internet. This digital nature has led to the popularity of standardized formats, ensuring compatibility across diverse platforms.

However, the ease of image creation and sharing, coupled with the vast amount of image data generated daily, may present unexpected challenges for storage and transmission. In 2024, an estimated 6 billion photos are taken daily, with 14 billion images shared on social media platforms. Especially in the context of thesis, the resulting volume of image data necessitates efficient delivery. Fortunately, optimization methods exist, relying on core concepts of image representation, encoding, and compression, to ensure efficient delivery across networks. We will explore these in the subsequent sections.

2.1.1 Digital imaging

Digital imaging, like traditional photography, relies on the fundamental principle of capturing light to create a visual representation. Most image acquisition methods, whether they involve cell phone cameras, DSLR (digital single-lens reflex) and mirrorless cameras, or even scanners, are variations of the classical optical camera. Concepts such as the relationship between aperture and shutter speed to achieve proper exposure are universal, regardless of whether the capture method is analog or digital.

However, digital imaging distinguishes itself by representing images as discrete, two-dimensional arrays of numerical data. This involves transforming light into a grid of distinct pixels. Each pixel may contain one or more channels — components that record different aspects of light, often corresponding to how our eyes perceive color. If only one channel is present, the image is effectively grayscale, as was the case with the first digital camera invented in 1975 by Steven Sasson.

To create a digital image, the camera must first capture the light that's focused onto its sensor. The position of the camera and its lens determine what portion of the scene is captured by each pixel on the sensor — this is known as spatial sampling. The camera also controls how long it gathers light at each pixel, called temporal sampling, and is determined by the exposure time, or using photographic terminology, shutter speed. Thanks to the photovoltaic effect, the amount of the light captured can be converted into digital value through a process called quantization. The number of distinct values that can be represented for each pixel depends on the bit depth of the image.

Digital cameras primarily use two image sensor technologies: CCD (Charge-Coupled Device) and CMOS (Complementary Metal-Oxide-Semiconductor). Early digital cameras utilized CCD sensors, which offer high sensitivity and low noise due to their charge transfer method, where all charge is shifted across the chip to a single output node. However, this sequential readout makes them slower for continuous shooting or video capture. CCDs also tend to consume more power and are more prone to blooming, where overexposed pixels overflow. In contrast, CMOS sensors have seen rapid development since the 1990s. While initially lagging in sensitivity and noise performance comparing to CCDs, CMOS sensors have greatly improved through advancements like on-chip noise reduction. Their lower cost, lower power consumption, and faster readout speeds have led to CMOS sensors becoming more popular than CCDs and being found in a wide range of devices, from smartphones and webcams to DSLRs, mirrorless cameras, and even scientific imaging equipment.

To accurately capture color, most digital cameras use a color filter array (CFA) placed over the sensor. The most common type is the Bayer filter, named after its inventor, Bryce Bayer, working at the time at Eastman Kodak. This filter arranges red, green, and blue filters in a specific pattern over the pixels. Because the human eye is more sensitive to green light, or to be specific, because *green signal contributes more to the luminance signal than the red or blue signals* [], in a typical Bayer filter we can find twice as many green filters as red or blue. Cameras capture a full-color image by interpolating the missing color information for each pixel, in a process known as demosaicing.

Following the Bayer filter example, it is common for color images to have three channels, such as red, green, and blue (RGB), each with its own bit depth. If each channel in an RGB image has an 8-bit resolution, it allows for 256 different values per channel, resulting in over 16 million possible color combinations for each pixel. The resolution of an image refers to the total number of pixels it contains and is often measured in megapixels (millions of pixels). Higher resolution images have more pixels, what allows for larger prints and more detailed displays. The image resolution, along with the bit depth, directly affects the size of the uncompressed image file. However, because images often contain areas of similar color, compression techniques, discussed in the next section, are frequently employed to reduce file sizes, sometimes even during the image capture process within the camera devices.

2.1.2 Data and image compression

Digital data, whether representing text, audio, or images, is fundamentally a collection of bits. Data compression techniques aim to represent the same information using fewer bits, which is important for addressing physical constraints of digital systems, such as bandwidth, storage density, and power consumption. In 1948, C. Shannon introduced the concept of entropy as a measure of information content. Entropy defines the fundamental limit for lossless data compression algorithms, providing a theoretical target for minimizing the number of bits needed to represent information. To effectively leverage data compression techniques, it's important to differentiate between lossless and lossy compression, recognizing the trade-offs between preserving data integrity and achieving higher compression ratios.

Lossless compression methods enable reversible transformations of data into more compact forms, while preserving all original information. While this thesis primarily focuses on image processing, we would like to note that lossless techniques have much broader applications and can be applied to virtually any form of digital data, including text and binary files. This is crucial for efficient data transmission of any kind, where lossless compression reduces bandwidth requirements, but allows recipients to access the original information in its unaltered form.

Run-length encoding is a simple technique that exploits repetitions in data. For example, a sequence of the same 20 consecutive characters *a* can be efficiently stored as *(20, a)* instead of storing each individual character. This works especially well if the set of unique symbols in the data is small, and the data contains long *runs* of the same symbol.

Huffman coding takes a statistical approach by constructing a binary tree (*Huffman tree*) based on symbol probabilities. Leafs representing higher probability symbols are positioned closer to the root, with shorter code lengths. This results in a variable-length code, where the length of a symbol's bit sequence is inversely proportional to its frequency in

the data. By assigning shorter codes to common symbols, Huffman coding optimizes code length, reducing the overall number of bits needed.

Arithmetic coding takes yet another approach, encoding the entire message as a single fraction between the values of 0 and 1. The exact value is determined by repeatedly subdividing the interval based on the probabilities of the symbols. Each symbol narrows the interval, and the final fraction precisely represents the entire sequence, often allowing it to be represented with fewer bits than the original message. However, this method requires the probability distribution of the symbols being encoded to be shared between the encoder and the decoder.

In addition to the methods described above that can help us understand the basics of compression, practical implementations of lossless compression often rely on dictionary-based techniques. The *Lempel-Ziv* algorithms, including LZ77, LZ78, and the *Lempel-Ziv-Welch* (LZW) algorithm, are important examples of this approach. LZ77 utilizes a sliding window to locate matches in preceding data, while LZ78 constructs a dictionary of previously encountered phrases. Combining LZ77 with Huffman coding, *Deflate* often achieves higher compression ratios than LZ77 alone.

Asymmetric Numeral Systems (ANS) is a relatively new approach to entropy coding that combines the speed of Huffman coding with the compression rate of arithmetic coding. Unlike arithmetic coding, which uses two fractional numbers to represent an interval, ANS encodes data directly as a single natural number. This simpler representation, using just one integer, reduces computational complexity and allows for higher compression and decompression throughput. We expect ANS to gain more popularity across various domains, including web delivery and image compression, as research and implementation efforts continue.

Lossy compression methods, unlike lossless techniques, offer significantly higher compression ratios by selectively discarding information. This trade-off is often acceptable for media files like images and audio, where human perception of minor data alterations may be limited, but may be not acceptable for other data types, such as computer programs.

Transform coding is a common technique used in lossy compression that serves two purposes. By compacting the original data samples into the transform coefficients with increasingly smaller variances, it is eventually able to remove smaller coefficients with negligible reconstruction errors. Additionally, decorrelated coefficients can be quantized and entropy coded individually without losing compression performance. By applying a transform, such as the *Discrete Cosine Transform* (DCT), we can separate the information into components that vary in importance with respect to human perception. In lossy image compression, transforms allow for the selective removal of high-frequency components which correspond to image details less perceptible to human eyes.

Subsampling exploits the characteristics of human perception by reducing the spatial resolution of specific components within the data. In image compression, this often involves relative downsampling the chrominance components (*chroma*), as the human eye is less sensitive to color detail compared to luminance (*luma*) detail. The reduction in chrominance resolution contributes to a smaller file size without significantly impacting the overall perceived image quality.

Evaluating the effectiveness of lossy compression involves measuring the perceived difference between the original and compressed images. Selected evaluation metrics include Peak Signal-to-Noise Ratio (PSNR), Structural Similarity Index Measure (SSIM), and

Butteraugli. While metrics like PSNR and SSIM provide objective measurement of the differences between pictures, the ultimate measure of lossy compression effectiveness will eventually depend on the subjective perception of its observers.

Another factor to consider with lossy compression is generation loss, which refers to the accumulated degradation of quality that occurs when a compressed file is repeatedly edited and recompressed. Each recompression cycle may introduce additional information loss, potentially leading to noticeable artifacts and a significant reduction in overall subjective quality of the image. In the next section, we will explore different image formats that utilize the lossy and lossless compression techniques described above.

2.1.3 Raster and vector graphics

Image formats are standardized ways of organizing digital image data and metadata within files. They ensure compatibility and allow different software, devices, and operating systems to work with the same images. Two fundamental categories of image formats are *raster* and *vector*.

Raster images are fundamentally composed of a grid of individual pixels, each representing a specific color. This structure directly corresponds to the way digital cameras capture images and is well-suited for representing photographs and other visuals with continuous tones. The term raster originates from the scanning pattern (*raster scan*) of CRT monitors, where the electron beam illuminates the screen pixel by pixel, line by line. In addition to color data, some raster formats incorporate an alpha channel, which controls each pixel's transparency. Image quality in raster formats is fundamentally connected to the resolution, measured by the total number of pixels. When scaling a raster image beyond its original size, preserving a fixed number of pixels can lead to a visible pixelation - a situation when individual pixels become visible to a naked eye. The alternative approach of introducing new pixels also has drawbacks: traditional *generating a raster image with a higher resolution than its source* using pixel interpolation methods may cause a visible loss of sharpness, also known as *edge blurring*.

Vector graphics offer a fundamentally different approach, representing visuals using geometrical primitives like lines, curves, circles, ellipses, or polylines and polygons. These primitives are defined by their mathematical properties, not as individual pixels. For instance, a circle may be defined by its center point coordinates and radius, while a curve might be represented using its end point and control points. This allows vector graphics to be inherently scalable, as the same geometrical definitions are used to render the graphic at any size.

While the following sections of our thesis will focus on raster image optimizations, a brief comparison of image optimization techniques for both vector and raster images is beneficial for the comprehensive understanding of the topic. Pixel-based methods, such as reducing image resolution, the bit depth, or applying lossy compression can not be directly translated to their vector counterparts, as vector graphics do not rely on a pixel grid. Instead, byte size of vector graphics may be optimized through methods unique to their geometrical nature. While lossless compression methods still apply, further optimizations are typically achieved by reducing the precision of values defining geometrical primitives, simplifying complex shapes by using fewer defining points, and, in formats supporting reusable styles, minimizing the number of distinct styles used.

2.1.4 Raster image formats

Building upon our understanding of raster graphics and the various compression techniques, we will now explore specific raster image formats. While often based on shared compression principles, different formats use different strategies for storing and compressing pixel data.

The existence of multiple formats is driven by the need to support a diverse range of viewers (clients), each with varying capabilities and limitations. To ensure optimal image delivery across this spectrum, it becomes necessary to provide assets in multiple formats, resolutions, and quality settings. This approach maximizes performance and visual fidelity for each user, but introduces the significant challenge of increased complexity in asset management.

Consequently, they vary in file size, image quality, and supported features. In the domain of image optimization, existence of different formats necessitates supporting multiple formats, and often multiple versions of the same image within a given format, to accommodate the varying capabilities of end-user devices and software.

BMP (Bitmap) is one of the oldest and simplest raster formats, developed by Microsoft for Windows. It primarily uses lossless compression through run-length encoding (RLE), or can store images uncompressed. While BMP supports various color depths, its compression is generally inefficient compared to more modern formats, leading to larger file sizes. BMP's simplicity makes it compatible with a wide range of software, but its lack of advanced features and efficient compression limits its use in scenarios where file size is a major concern. BMP is better suited for situations requiring simple, uncompressed image storage or compatibility with older systems.

GIF (Graphics Interchange Format), introduced in 1987, was an early internet staple due to its support for animation and lossless compression based on the LZW algorithm. However, GIF is limited to an 8-bit palette, meaning it can only display 256 colors. This restriction makes it unsuitable for photographs but still viable for graphics with limited color ranges, such as logos or simple icons. GIF also supports transparency, allowing for irregularly shaped, non-rectangular images. While GIF has been largely superseded by PNG for static images, its animation capabilities continue to make it relevant for short, looping animations, especially in online contexts.

PNG (Portable Network Graphics) was developed as a patent-free alternative to GIF, offering improved lossless compression and extended features. Using the Deflate algorithm, PNG achieves better compression than GIF without sacrificing image quality. It supports a wider range of color depths, including 24-bit and 48-bit, making it suitable for both simple graphics and photographs. PNG also offers full alpha channel support, allowing for smooth transparency gradients. These features have made PNG the preferred choice for web graphics requiring high quality and transparency, such as logos, icons, and interface elements.

JPEG (Joint Photographic Experts Group) is the dominant format for photographic images on the web and in digital cameras. JPEG employs lossy compression based on the Discrete Cosine Transform (DCT), quantization, and entropy coding, achieving significant file size reductions. Users can typically adjust the compression level, allowing for a trade-off between image quality and file size. While JPEG excels at compressing continuous-tone images like photographs, the lossy nature of the compression can introduce visible arti-

facts, especially around sharp edges and areas of high contrast. JPEG does not support transparency, making it unsuitable for graphics that require it.

JPEG 2000 was designed as a successor to JPEG, offering improved compression efficiency and a wider range of features. Instead of DCT, JPEG 2000 utilizes wavelet transforms, which provide better image quality at higher compression ratios. It supports both lossy and lossless compression, as well as advanced features like region of interest (ROI) encoding, allowing specific areas of an image to be encoded with higher quality than others. While JPEG 2000 offers technical advantages, its adoption has been limited due to factors like licensing costs and the widespread use of the original JPEG format. It finds niche applications in fields like medical imaging and digital archiving, where its advanced features are beneficial.

WebP is a relatively new format developed by Google, specifically for web use. It supports both lossy and lossless compression, providing smaller file sizes compared to JPEG and PNG, respectively. WebP's lossy compression is based on the VP8 video codec, while its lossless mode uses a variety of techniques, including dictionary coding, Huffman coding, and color indexing. WebP also supports transparency and animation, potentially replacing both PNG and GIF in many web applications. While browser support for WebP has been steadily increasing, it is still not universally supported at the time of writing.

AVIF (AV1 Image File Format) is a modern image format derived from the AV1 video codec, designed to offer superior compression compared to existing formats like JPEG, PNG, and even WebP. AVIF supports both lossy and lossless compression and boasts a wide range of features, including support for high dynamic range (HDR) and wide color gamut (WCG) images, as well as transparency. AVIF generally offers significant file size reductions compared to JPEG while maintaining comparable visual quality. Its support for advanced features like HDR and WCG makes it suitable for high-quality images and modern display technologies.

JPEG XL is another modern image format designed to be a versatile successor to JPEG. It offers both lossy and lossless compression, outperforming JPEG in terms of compression efficiency. One of JPEG XL's key strengths is its ability to losslessly recompress existing JPEG files, reducing their size without any further quality loss. This feature, along with progressive decoding capabilities, can improve user experience, particularly for those on slower networks. JPEG XL also supports animation, alpha channels, and a wide range of color depths, including HDR, making it suitable for diverse use cases.

2.2 Cloud Computing

2.2.1 Cloud deployment models

2.2.2 Event-driven architectures

2.2.3 Cloud service benchmarking

2.2.4 Performance variability of serverless platforms

2.3 Image Optimization for the Web

2.3.1 Impact of images on browsing experience

2.3.2 Bandwidth usage, visual quality, and browser support evolution

2.3.3 Optimizing latency of image delivery

2.3.4 Deployment methods for cloud-based image optimization

3 System Design

3.1 Image processing workload design

3.1.1 Serverless environment selection

3.1.2 Image manipulation library selection

3.1.3 Input image and output image selection

3.2 Benchmark design

3.2.1 Serverless platform and underlying hardware

3.2.2 Time of execution

3.2.3 Output formats and image widths

3.2.4 Available CPU / memory

3.2.5 Cold starts

3.3 Workload generation and benchmark execution

3.3.1 Serverless benchmark deployment

3.3.2 Scheduling workload executions

3.3.3 Results collection and data analysis

4 Results

4.1 Performance metrics

4.1.1 Cold start duration

4.1.2 Image processing duration

4.1.3 Memory utilization

4.2 Performance variability

4.2.1 General variability in performance metrics

4.2.2 Temporal variability

4.2.3 By image format

4.2.4 By platform configuration

5 Discussion

5.1 Performance variability analysis

5.2 User vs operator perspective

5.3 Cloud computing as a utility

6 Related Work

6.1 Serverless performance research

6.2 Image processing research

7 Conclusion