

# Practical FaaS Performance Variability Implications in the Serverless Image Optimization Scenario

Master's Thesis

## Author

Piotr Witkowski

414360

witkowski@campus.tu-berlin.de

## Advisor

Trever Schirmer

## Examiners

Prof. Dr.-Ing. David Bermbach

Prof. Dr. habil. Odej Kao

**Technische Universität Berlin, 2024**

Fakultät Elektrotechnik und Informatik

Fachgebiet Scalable Software Systems

# Practical FaaS Performance Variability Implications in the Serverless Image Optimization Scenario

Master's Thesis

Submitted by:  
Piotr Witkowski  
414360  
witkowski@campus.tu-berlin.de

Technische Universität Berlin  
Fakultät Elektrotechnik und Informatik  
Fachgebiet Scalable Software Systems

2024

Hiermit versichere ich, dass ich die vorliegende Arbeit eigenständig ohne Hilfe Dritter und ausschließlich unter Verwendung der aufgeführten Quellen und Hilfsmittel angefertigt habe. Alle Stellen die den benutzten Quellen und Hilfsmitteln unverändert oder sinngemäß entnommen sind, habe ich als solche kenntlich gemacht.

Sofern generische KI-Tools verwendet wurden, habe ich Produktnamen, Hersteller, die jeweils verwendete Softwareversion und die jeweiligen Einsatzzwecke (z. B. sprachliche Überprüfung und Verbesserung der Texte, systematische Recherche) benannt. Ich verantworte die Auswahl, die Übernahme und sämtliche Ergebnisse des von mir verwendeten KI-generierten Outputs vollumfänglich selbst.

Die Satzung zur Sicherung guter wissenschaftlicher Praxis an der TU Berlin vom 8. März 2017\* habe ich zur Kenntnis genommen.

Ich erkläre weiterhin, dass ich die Arbeit in gleicher oder ähnlicher Form noch keiner anderen Prüfungsbehörde vorgelegt habe.

---

(Unterschrift) Piotr Witkowski, Berlin, 2. Januar 2025

---

\*[https://www.static.tu.berlin/fileadmin/www/10000060/FSC/Promotion\\_\\_Habilitation/Dokumente/Grundsatzgute\\_wissenschaftliche\\_Praxis\\_2017.pdf](https://www.static.tu.berlin/fileadmin/www/10000060/FSC/Promotion__Habilitation/Dokumente/Grundsatzgute_wissenschaftliche_Praxis_2017.pdf)

## Abstract

## Kurzfassung

# Contents

|          |                                                                          |           |
|----------|--------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                                                      | <b>8</b>  |
| 1.1      | Objectives of the Thesis . . . . .                                       | 8         |
| 1.2      | Structure . . . . .                                                      | 9         |
| <b>2</b> | <b>Background</b>                                                        | <b>10</b> |
| 2.1      | Imaging and Photography . . . . .                                        | 10        |
| 2.1.1    | Digital imaging and image processing . . . . .                           | 11        |
| 2.1.2    | Raster and vector graphics . . . . .                                     | 11        |
| 2.1.3    | Image data compression . . . . .                                         | 11        |
| 2.1.4    | Image formats and their usage . . . . .                                  | 11        |
| 2.2      | Cloud Computing . . . . .                                                | 11        |
| 2.2.1    | Cloud deployment models . . . . .                                        | 11        |
| 2.2.2    | Event-driven architectures . . . . .                                     | 11        |
| 2.2.3    | Cloud service benchmarking . . . . .                                     | 11        |
| 2.2.4    | Performance variability of serverless platforms . . . . .                | 11        |
| 2.3      | Image Optimization for the Web . . . . .                                 | 11        |
| 2.3.1    | Impact of images on browsing experience . . . . .                        | 11        |
| 2.3.2    | Bandwidth usage, visual quality, and browser support evolution . . . . . | 11        |
| 2.3.3    | Optimizing latency of image delivery . . . . .                           | 11        |
| 2.3.4    | Deployment methods for cloud-based image optimization . . . . .          | 11        |
| <b>3</b> | <b>System Design</b>                                                     | <b>12</b> |
| 3.1      | Image processing workload design . . . . .                               | 12        |
| 3.1.1    | Serverless environment selection . . . . .                               | 12        |
| 3.1.2    | Image manipulation library selection . . . . .                           | 12        |
| 3.1.3    | Input image and output image selection . . . . .                         | 12        |
| 3.2      | Benchmark design . . . . .                                               | 12        |
| 3.2.1    | Serverless platform and underlying hardware . . . . .                    | 12        |
| 3.2.2    | Time of execution . . . . .                                              | 12        |
| 3.2.3    | Output formats and image widths . . . . .                                | 12        |
| 3.2.4    | Available CPU / memory . . . . .                                         | 12        |
| 3.2.5    | Cold starts . . . . .                                                    | 12        |
| 3.3      | Workload generation and benchmark execution . . . . .                    | 12        |
| 3.3.1    | Serverless benchmark deployment . . . . .                                | 12        |
| 3.3.2    | Scheduling workload executions . . . . .                                 | 12        |
| 3.3.3    | Results collection and data analysis . . . . .                           | 12        |
| <b>4</b> | <b>Results</b>                                                           | <b>13</b> |
| 4.1      | Performance metrics . . . . .                                            | 13        |
| 4.1.1    | Cold start duration . . . . .                                            | 13        |
| 4.1.2    | Image processing duration . . . . .                                      | 13        |
| 4.1.3    | Memory utilization . . . . .                                             | 13        |
| 4.2      | Performance variability . . . . .                                        | 13        |
| 4.2.1    | General variability in performance metrics . . . . .                     | 13        |
| 4.2.2    | Temporal variability . . . . .                                           | 13        |
| 4.2.3    | By image format . . . . .                                                | 13        |
| 4.2.4    | By platform configuration . . . . .                                      | 13        |

|          |                                            |           |
|----------|--------------------------------------------|-----------|
| <b>5</b> | <b>Evaluation</b>                          | <b>14</b> |
| 5.1      | Performance variability analysis . . . . . | 14        |
| 5.2      | User vs operator perspective . . . . .     | 14        |
| 5.3      | Cloud computing as a utility . . . . .     | 14        |
| <b>6</b> | <b>Related Work</b>                        | <b>15</b> |
| 6.1      | Serverless performance research . . . . .  | 15        |
| 6.2      | Image processing research . . . . .        | 15        |
| <b>7</b> | <b>Conclusion</b>                          | <b>16</b> |

# 1 Introduction

The digital revolution has fundamentally transformed the way we create, consume, and interact with visual information. Images are no longer only physical items around us, they are inherent components of our online experience, shaping perceptions and influencing decisions. As the Internet continues to evolve, so too grows the demand for rich and engaging visual content. This demand, however, comes at a cost. Images, while visually appealing, are significantly impacting website performance, not only by generating the highest number of requests, but also by consuming the most bytes during page load of an average site. Existing studies have shown a strong correlation between page load times and user engagement, conversion rates, and even search engine rankings. Minimizing image size is therefore fundamental for efficient bandwidth use and plays a crucial role in optimizing website performance.

Image optimization practices for the web have evolved over time, following developments of the state-of-the-art image formats and codecs. These solutions perform operations such as image resizing and format conversion, and may allow for storing optimized versions of original files for future use. While the optimization itself is a relatively specialized task, its deployment models follow familiar and established patterns in the general software industry: we can find them in software-as-a-service offerings, integrated into web hosting and content delivery platforms, but also as open-source solutions for self-hosting.

This last option, leveraging open-source image processing libraries, presents a compelling opportunity for our research. Unlike proprietary solutions, open-source projects offer the flexibility to be deployed across a variety of environments, granting builders granular control over the optimization process. This allows not only to measure influence of different libraries and optimization parameters, but also – and crucially for this thesis – enables realistic benchmarking of the underlying platforms on which these operations are performed. By employing these workloads, our research aims to accurately measure cloud service performance under the realistic conditions of real-world image optimization scenarios.

Given the event-driven nature of the optimization workload in response to requests for images, serverless cloud compute emerges as a well-suited deployment option for these systems. However, the varying demands for image processing, including different image sizes and diverse range of image formats present on the web, raise questions about the suitability and performance consistency of these platforms for the use-case. Furthermore, the emergence of function as a service (FaaS) platforms introduces a new dimension to cloud-based, self-hosted image optimization. FaaS offers a compelling execution environment for image processing tasks thanks to its on-demand scalability and event-driven nature. However, the performance variability inherent in FaaS platforms, attributed to diverse factors, may impact efficiency of the process and applicability of the environment for the use case.

## 1.1 Objectives of the Thesis

This thesis aims to investigate whether FaaS performance variability has measurable and meaningful influence for image optimization workloads. We will define set of representative metrics related to image optimization, reflecting realistic web usage. We will design and implement the optimization workflows for selected serverless environments. We will design and implement scheduling mechanism allowing us to simulate requests for optimized images. We will collect performance metrics across all parameters, apply statistical



methods for result comparisons, and visualize the results. We will investigate if more intensive optimization can be cost-efficient, considering optimization result caching, and content reuse for multiple viewers. We will provide recommendations that are within the scope of the benchmark, and answer the primary research question in the context of the scenario.

## **1.2 Structure**

## 2 Background

This chapter provides readers with foundations for understanding the challenges and opportunities of serverless image processing. We begin by revisiting fundamental principles of imaging and photography. Subsequently, we describe core digital imaging concepts, emphasizing aspects relevant to web performance. We then examine the characteristics of serverless cloud computing, focusing on factors contributing to performance variability. Finally, we explore image optimization for the web, reviewing established techniques and the landscape of cloud-based solutions.

### 2.1 Imaging and Photography

*Imaging* is a broad term describing any process that helps us in visualizing and understanding physical objects, abstract concepts, or even sets of data. While this may be a purely philosophical activity, for the purpose of this thesis, we will define imaging as creating a visual representation - *an image* - of physical objects and scenes. Humans are able to perceive and interpret these images thanks to our visual system, that has evolved over millions of years. Developed from the simple light-sensitive cells of early organisms to the sophisticated eyes of modern humans, this system is providing us with the ability to see color, depth, or movement.

Over centuries, scholars and artists showed great interest in this process of creating visual representations to understand our own senses, as well as to be able to depict the world around us more accurately. The natural phenomenon of light passing through a small aperture, called camera obscura, helped us understand fundamental principles of optics and image formation, and according to some, greatly contributed to realism of European art during Middle Ages. We've been refining tools and techniques to enhance our vision, and improve everyday life. Devices like telescopes and microscopes allowed us to observe objects at different scales, unveiling worlds that would otherwise remain invisible to the naked eye. Eyeglasses and contact lenses correct refractive errors, demonstrating that our visual experiences can not only be externally altered, but also subjectively improved, for millions of people.

Photography emerged in the 19th century as a groundbreaking invention, offering a novel approach to image creation. Unlike paintings or drawings, which rely on an artist's interpretation, photography directly captures and permanently records the physical world using light-sensitive materials. Early photographic methods required long exposure times, often making it possible to photograph only still objects and posed portraits. The daguerreotype, which used a silver-coated copper plate, was one of the first widely adopted photographic processes. Later, more versatile film technologies significantly reduced exposure times and made photography more accessible to the wider public. This ability to easily and accurately record the physical world had significant consequences. Newspapers, now able to include actual photographs, found a powerful new way to share news and social commentary through photojournalism. On a personal level, photography allowed people to preserve their own histories through portraits and the documentation of special occasions.

Since the beginning of the 21st century, digital photography has become the dominant form of capturing and sharing images. Initially limited to standalone devices, digital cameras are now integrated into smartphones and personal computers. Electronic sensors capture and convert light into digital data, facilitating easy manipulation and image sharing, including

over the Internet. This digital nature has led to the popularity of standardized formats, ensuring compatibility across diverse platforms.

However, the ease of image creation and sharing, coupled with the vast amount of image data generated daily, may present unexpected challenges for storage and transmission. In 2024, an estimated 6 billion photos are taken daily, with 14 billion images shared on social media platforms. Especially in the context of thesis, the resulting volume of image data necessitates efficient delivery. Fortunately, optimization methods exist, relying on core concepts of image representation, encoding, and compression, to ensure efficient delivery across networks. We will explore these in the subsequent sections.

### **2.1.1 Digital imaging and image processing**

### **2.1.2 Raster and vector graphics**

### **2.1.3 Image data compression**

### **2.1.4 Image formats and their usage**

## **2.2 Cloud Computing**

### **2.2.1 Cloud deployment models**

### **2.2.2 Event-driven architectures**

### **2.2.3 Cloud service benchmarking**

### **2.2.4 Performance variability of serverless platforms**

## **2.3 Image Optimization for the Web**

### **2.3.1 Impact of images on browsing experience**

### **2.3.2 Bandwidth usage, visual quality, and browser support evolution**

### **2.3.3 Optimizing latency of image delivery**

### **2.3.4 Deployment methods for cloud-based image optimization**

## 3 System Design

### 3.1 Image processing workload design

#### 3.1.1 Serverless environment selection

#### 3.1.2 Image manipulation library selection

#### 3.1.3 Input image and output image selection

### 3.2 Benchmark design

#### 3.2.1 Serverless platform and underlying hardware

#### 3.2.2 Time of execution

#### 3.2.3 Output formats and image widths

#### 3.2.4 Available CPU / memory

#### 3.2.5 Cold starts

### 3.3 Workload generation and benchmark execution

#### 3.3.1 Serverless benchmark deployment

#### 3.3.2 Scheduling workload executions

#### 3.3.3 Results collection and data analysis

## 4 Results

### 4.1 Performance metrics

#### 4.1.1 Cold start duration

#### 4.1.2 Image processing duration

#### 4.1.3 Memory utilization

### 4.2 Performance variability

#### 4.2.1 General variability in performance metrics

#### 4.2.2 Temporal variability

#### 4.2.3 By image format

#### 4.2.4 By platform configuration

## **5 Evaluation**

### **5.1 Performance variability analysis**

### **5.2 User vs operator perspective**

### **5.3 Cloud computing as a utility**

## **6 Related Work**

### **6.1 Serverless performance research**

### **6.2 Image processing research**

## 7 Conclusion